

MINSU KIM

Postdoctoral Fellow
Mila - Quebec AI Institute

✉ minsu.kim@mila.quebec

G google scholar

🌐 homepage

HIGHLIGHTS

- 40+ peer-reviewed papers, incl. 23 NeurIPS/ICML/ICLR (10 lead); 14 lead overall
- Joint projects with LLNL, Intel, Samsung, LG, KT, and LawZero, bridging research and industry applications
- Research areas:
 - **ML fundamentals**: exploration, credit assignment, probabilistic inference, generative modeling, distributed RL
 - **RL for post-training**: aligning LLMs, diffusion, GANs, VAEs, protein folding
 - **Trustworthy reasoning of LLMs**: red-teaming, uncertainty quantification, self-verification
 - **AI4Opt**: neural solvers for NP-hard problems, RL for multi-robot distributed scheduling
 - **AI4Science**: generative modeling and RL for molecules, 3D conformations, biological sequences
 - **AI4Sys**: ML-driven optimization for physical systems, including HBM, interconnects, and power distribution networks

EMPLOYMENT AND COLLABORATIONS

- 6/2025 – Present **Postdoctoral Fellow (Host: Prof. Yoshua Bengio)** **Mila**
- Trustworthy reasoning in LLM
 - *Joint affiliation with KAIST via the KAIST–Mila Prefrontal Research Collaboration (Profs. Sungjin Ahn and Sungsoo Ahn)*
- 6/2025 – Present **Academic Collaborator** **LawZero**
- Led collaboration on trustworthy LLMs; focusing on red-teaming and safe reasoning
 - Contributed to the uncertainty measurement of LLMs project with joint experiments and publications
- 1/2025 – 6/2025 **Research Intern (Supervisor: Prof. Yoshua Bengio)** **Mila**
12/2023 – 5/2024
- GFlowNets for LLM/LMM fine-tuning
 - GFlowNets for biological and chemical discovery

EDUCATION

- 3/2022 – 2/2025 **Ph.D. in Industrial and Systems Engineering (Advisor: Prof. Jinkyoo Park)** **KAIST**
Thesis: Off-policy training methods for probabilistic agents in combinatorial space
Presidential Best Ph.D. Thesis Award
- 3/2020 – 2/2022 **M.S. in Electrical Engineering (Advisor: Prof. Joung-ho Kim)** **KAIST**
Research: Applied reinforcement learning to routing problems in high-speed interconnection design
- 3/2015 – 2/2020 **B.S. in Mathematics and Computer Science (Dual Major)** **KAIST**

INDUSTRY-ACADEMIA COLLABORATION PROJECTS

- 2025 – Present **Scalable Multi-agent Off-policy Distributed Training of GFlowNets** **Intel–Mila**
- Research on distributed training pipelines and analyzing scaling laws on multi-GPU clusters
- 2024 – 2025 **Distributed Off-Policy RL for Efficient LLM Post-Training** **Lawrence Livermore National Lab (LLNL) - Mila**
- Co-authored work (NeurIPS 2025) on asynchronous off-policy RL to improve LLM fine-tuning efficiency
 - Achieved 4× faster wall-clock time in LLM post-training via asynchronous off-policy RL, enabling scalable deployment
- 2022 – 2023 **Multi-Robot Distributed Scheduling with Graph RL** **KT - KAIST**
- Cluster-robot scheduling using reinforcement learning on graphs
 - Led research publication of the project (AAAI 2024)
- 2020 – 2021 **AI for Semiconductor Design Automation considering Signal & Power Integrity** **Samsung-KAIST**
- Applied ML in high-speed interconnect design
 - Led research publication of the project (EPEPS 2022)

TOP 10 SELECTED PUBLICATIONS

(see THIS LINK for full publication)

Preprint, 2025	Active Attacks: Redteaming LLMs via Adaptive Environments Taeyoung Yun, Pierre-Luc St-Charles, Jinkyoo Park, Yoshua Bengio, and Minsu Kim
Preprint, 2025	Latent Veracity Inference for Identifying Errors in Stepwise Reasoning { Minsu Kim *, Jean-Pierre Falet*}, Oliver E. Richardson, Xiaoyin Chen, Moksh Jain, Sungjin Ahn, Sungsoo Ahn, and Yoshua Bengio
NeurIPS, 2025	On Scalable and Efficient Training of Diffusion Samplers Minkyu Kim, Kiyoung Seong, Dongyeop Woo, Sungsoo Ahn, and Minsu Kim
ICLR, 2025	Adaptive Teachers for Amortized Samplers { Minsu Kim *, Sanghyeok Choi*}, Taeyoung Yun, Emmanuel Bengio, Leo Feng, Jarrod Rector-Brooks, Sungsoo Ahn, Jinkyoo Park, Nikolay Malkin, and Yoshua Bengio
NeurIPS, 2024	Amortizing Intractable Inference in Diffusion Models for Vision, Language, and Control {Siddharth Venkatraman*, Moksh Jain*, Luca Scimeca*, Minsu Kim *, Marcin Sendera*} et al.
ICML 2024	Learning to Scale Logits for Temperature-Conditional GFlowNets { Minsu Kim *, Joohwan Ko*, Taeyoung Yun*}, Dinghui Zhang, Ling Pan, Taeyoung Yun, Woosung Kim, Jinkyoo Park, Emmanuel Bengio, and Yoshua Bengio
ICLR, 2024	Local Search GFlowNets Minsu Kim , Taeyoung Yun, Emmanuel Bengio, Dinghui Zhang, Yoshua Bengio, Sungsoo Ahn, and Jinkyoo Park
NeurIPS, 2023	Bootstrapped Training of Score-Conditioned Generator for Offline Design of Biological Sequences Minsu Kim , Federico Berto, Sungsoo Ahn, and Jinkyoo Park
NeurIPS, 2022	Sym-nco: Leveraging Symmetricity for Neural Combinatorial Optimization Minsu Kim , Junyoung Park, and Jinkyoo Park
NeurIPS, 2021	Learning Collaborative Policies to Solve NP-hard Routing Problems Minsu Kim , Jinkyoo Park, and Joungho Kim

SELECTED HONORS AND AWARDS (LINK: FULL RECORDS)

6/2025	CIFAR AI Safety Postdoctoral Fellowship (via Yoshua Bengio) Awarded by the Canadian AI Safety Institute at CIFAR, under the mentorship of Prof. Yoshua Bengio	CIFAR / Mila
4/2025	Jang Yeong Sil Fellowship Award A prestigious award recognizing early-career postdoctoral researchers with outstanding potential	KAIST
12/2023	Qualcomm Innovative Fellowship Award Paper: Sym-NCO: Leveraging Symmetricity for Neural Combinatorial Optimization	Qualcomm Korea
2021	Best Paper Award (First Author) Paper: Neural Language Model Enables Extremely Fast and Robust Routing on Interposer	DesignCon 2021

INVITED TALKS AND INDUSTRY PRESENTATIONS

2025	Rising Star Talk: Trustworthy Foundation Model	MBZUAI (Abu Dhabi)
2024	Invited Research Talk	Valence Labs (Virtual)
2025	Research Spotlight, LG Tech Day	LG AI Research (Seoul)
2023	Research Spotlight, Samsung AI Forum	Samsung Advanced Institute of Technology (Suwon)